

INDRA REDDY

AI/ML Engineer

San Francisco Bay area, CA, United States | +1 (469) 224-0431 | indrareddyds009@gmail.com | [LinkedIn](#) | [Portfolio](#)

SUMMARY

LLM Engineer with 4+ years building production-grade GenAI and agentic AI systems in finance and SaaS. Hands-on experience across the LLM application lifecycle: prompt and context engineering, retrieval-augmented generation (RAG), fine-tuning, model evaluation, and inference optimization on Kubernetes/Triton. Strong MLOps and data engineering foundation across AWS, GCP, and Azure

SKILLS

Languages & Frameworks: Python (Pandas, NumPy, PyTorch, TensorFlow, Hugging Face), SQL, Bash
ML & AI: LLM fine-Tuning, RAG pipelines, NLP, Reinforcement learning, Anomaly Detection, XGBoost, Agentic AI.
MLOps & CI/CD: MLflow, Kubeflow, Airflow, AWS SageMaker, Jenkins, Drift Monitoring, Automated Retraining
Data Engineering: Apache Kafka, Spark Streaming, Snowflake, dbt, Snowpark, Real-Time ETL
Cloud & Infra: AWS (S3, Redshift, EKS, SageMaker, DGX Cloud), GCP, Azure, Kubernetes, Docker, Triton
Vector / Search: Pinecone, FAISS, Milvus, Semantic Search, Embedding-based Retrieval
Monitoring: Prometheus, Grafana, ELK Stack, Custom Model Evaluation Dashboards

EDUCATION

Master's in computer science | Trine University, Detroit, Michigan, USA

Bachelor of Technology | AKS University, Satna, Madhya Pradesh, India

WORK EXPERIENCE

GEICO, USA | AI Engineer

Aug 2026 - Present

- Building/enhancing ML models for auto insurance underwriting and risk scoring, using structured policyholder and claims data
- Developing pricing/rating models to improve premium accuracy and risk segmentation across policy lines
- Partnering with actuarial and underwriting teams to translate business rules into ML features and model constraints
- Monitoring model performance and fairness/bias metrics for risk models in production

State Street, USA | AI/ML Engineer | Redwood city, CA

Sep 2024 - Aug 2026

- Designed and maintained ML pipelines processing 50M+ daily financial transactions across risk, compliance, and reporting workflows.
- Fine-tuned large language models for financial document understanding, improving extraction accuracy.
- Built RAG-based retrieval systems for compliance Q&A, cutting manual analyst review time in a high-volume regulatory environment.
- Developed real-time anomaly detection models using Kafka & XGBoost contributed to a fraud prevention system flagging \$12M+ in suspicious transactions annually.
- Reduced model inference latency by 40% through quantization, batching, and Triton Inference Server optimizations.
- Implemented automated retraining pipelines (MLflow & Airflow) to handle distributional shift in streaming financial data.
- Deployed containerized inference services on EKS and monitored via Prometheus & Grafana dashboards.
- Improved risk scoring model precision through iterative feature engineering on structured transaction data.

Freshworks Inc, India | Data Scientist | Chennai, India

Mar 2020 - Nov 2022

- Built customer churn prediction models (XGBoost & logistic regression) improving identification accuracy by 29% across SaaS product lines.
- Designed collaborative-filtering recommendation engine that increased in-app engagement within the customer support platform.
- Optimized SQL pipelines and data transformation layers, reducing reporting query latency by 38% across business intelligence dashboards.
- Conducted statistical analysis uncovering \$2.5M in annual revenue leakage from billing and contract edge cases.
- Applied NLP (BERT-based classifiers) to automate support ticket classification, improving routing efficiency.
- Built A/B testing framework enabling faster experimentation cycles and evidence-based product decisions.
- Automated reporting pipelines saving 400+ operational hours per quarter; delivered stakeholder dashboards in Tableau.
- Developed fraud detection rules and ML signals that reduced false-positive rate and improved customer experience.

CERTIFICATIONS

AWS Certified Solutions Architect

PROJECTS

Autonomous Customer Support Agent (Agentic AI + RAG)

- Built an AI-powered customer support agent combining RAG and Agentic AI capabilities.
- Enabled the agent to retrieve relevant knowledge base data and take actions like ticket creation and status updates.
- Integrated CRM and ticketing systems via REST APIs for real-world task execution.
- Implemented conversation memory and context tracking for better user interactions.
- Reduced average response time by 70% and improved customer satisfaction.